

RIP – routing z użyciem protokołu dynamicznego

Routing Information Protocol (RIP)

- W ramach zajęć zostanie przeprowadzona analiza sieci, składającej się z pięciu ruterów wirtualnych.
- Adresy IP poszczególnych interfejsów tych urządzeń zostały już skonfigurowane.
- W ramach zajęć należy:
 - Zapoznać się z oprogramowaniem **quagga**, który jest dostępny w systemie operacyjnym Linux i zapewnia implementację protokołów routingu opartych na TCP/IP:
<http://www.nongnu.org/quagga/docs/quagga.html>
 - Skonfigurować protokół RIP z użyciem polecenia **telnet**
 - Dokonać analizy komunikatów protokołu RIP z użyciem polecenia **tcpdump**
 - Dokonać propagacji trasy domyślnej w ramach protokołu RIP
 - Zaobserwować dynamikę działania protokołu RIP

RIP – ruting z użyciem protokołu dynamicznego *Routing Information Protocol (RIP)*

■ **Głównym celem zajęć jest**

- konfiguracja dynamicznego protokołu routingu RIP opartego na algorytmie wektora odległości (ang. *distance vector*)
- zrozumienie zasad i mechanizmów działania protokołu RIP
- lepsze zapoznanie się ze środowiskiem **kathara**

kathara lab


rip

Version	1.2
Author(s)	G. Di Battista, M. Patrignani, M. Pizzonia, F. Ricci, M. Rimondini
E-mail	contact@kathara.org
Web	http://www.kathara.org/
Description	experiences with the ripv2 distance vector routing protocol – kathara version of the rip lab of netkit (vers. 2.4)

copyright notice

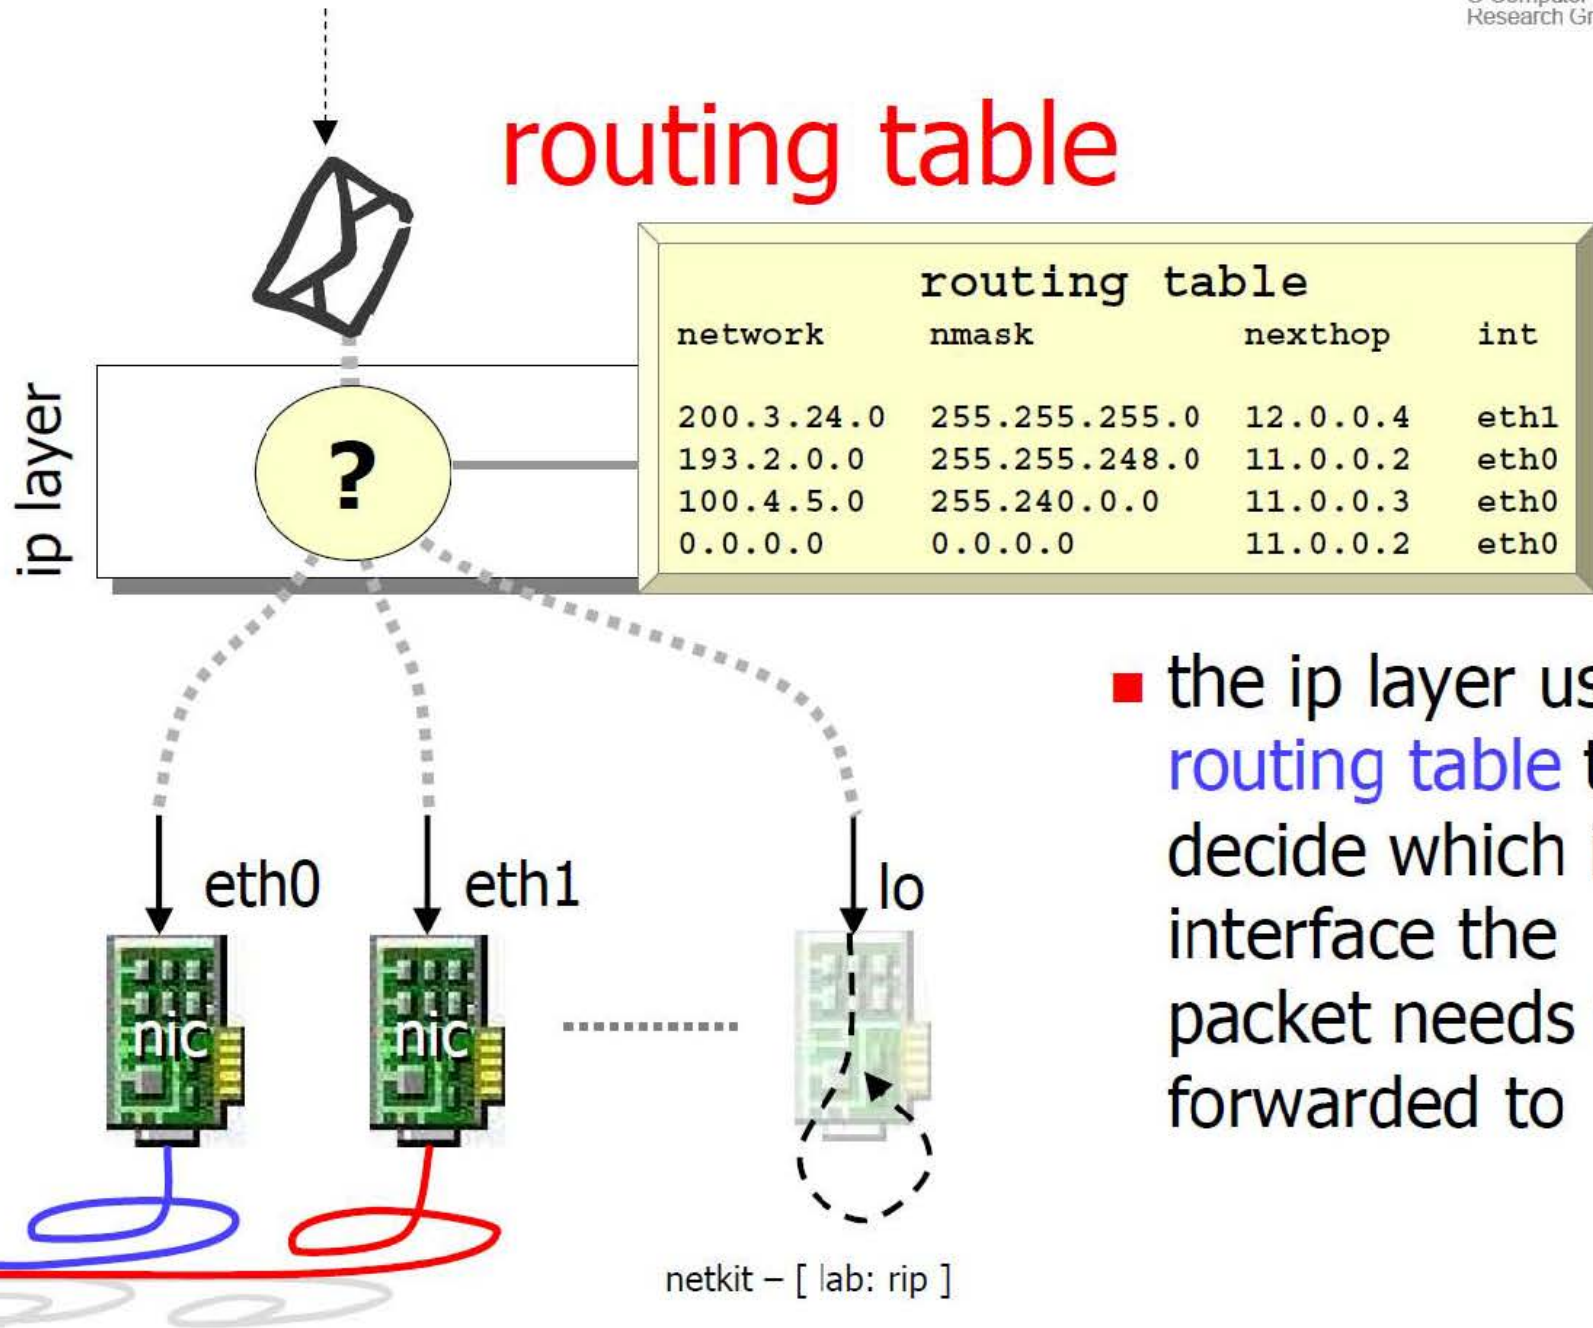
- All the pages/slides in this presentation, including but not limited to, images, photos, animations, videos, sounds, music, and text (hereby referred to as “material”) are protected by copyright.
- This material, with the exception of some multimedia elements licensed by other organizations, is property of the authors and/or organizations appearing in the first slide.
- This material, or its parts, can be reproduced and used for didactical purposes within universities and schools, provided that this happens for non-profit purposes.
- Information contained in this material cannot be used within network design projects or other products of any kind.
- Any other use is prohibited, unless explicitly authorized by the authors on the basis of an explicit agreement.
- The authors assume no responsibility about this material and provide this material “as is”, with no implicit or explicit warranty about the correctness and completeness of its contents, which may be subject to changes.
- This copyright notice must always be redistributed together with the material, or its portions.

hosts need routing




- each host with a network stack performs some elementary routing
- at the very least, the network stack may be used to access local services (e.g., XWindows)
- the host must decide when a packet needs to be sent to the **network interface card (nic)** and when it needs to be bounced to the **loopback interface (lo)**

routing table



- the ip layer uses a **routing table** to decide which is the interface the packet needs to be forwarded to




- a **router** (also called **gateway** or **intermediate-system**)
 - has more than one network interface card
 - feeds incoming ip packets (that are not for the router itself) back in the routing process
 - this operation is called **relaying** or **forwarding**


routing protocols

- routing protocols are used to automatically update the routing tables
- they fall into two main categories:
 - link-state routing protocols
 - approach: send the minimum information to everyone
 - each router reconstructs the whole network graph and computes a shortest path tree to all destinations
 - examples: is-is, ospf
 - distance-vector routing protocols
 - approach: send all your information to a few
 - update your routing information based on what you hear
 - examples: rip, bgp
- in this lab we will see an example of **RIPv2** protocol on quagga boxes

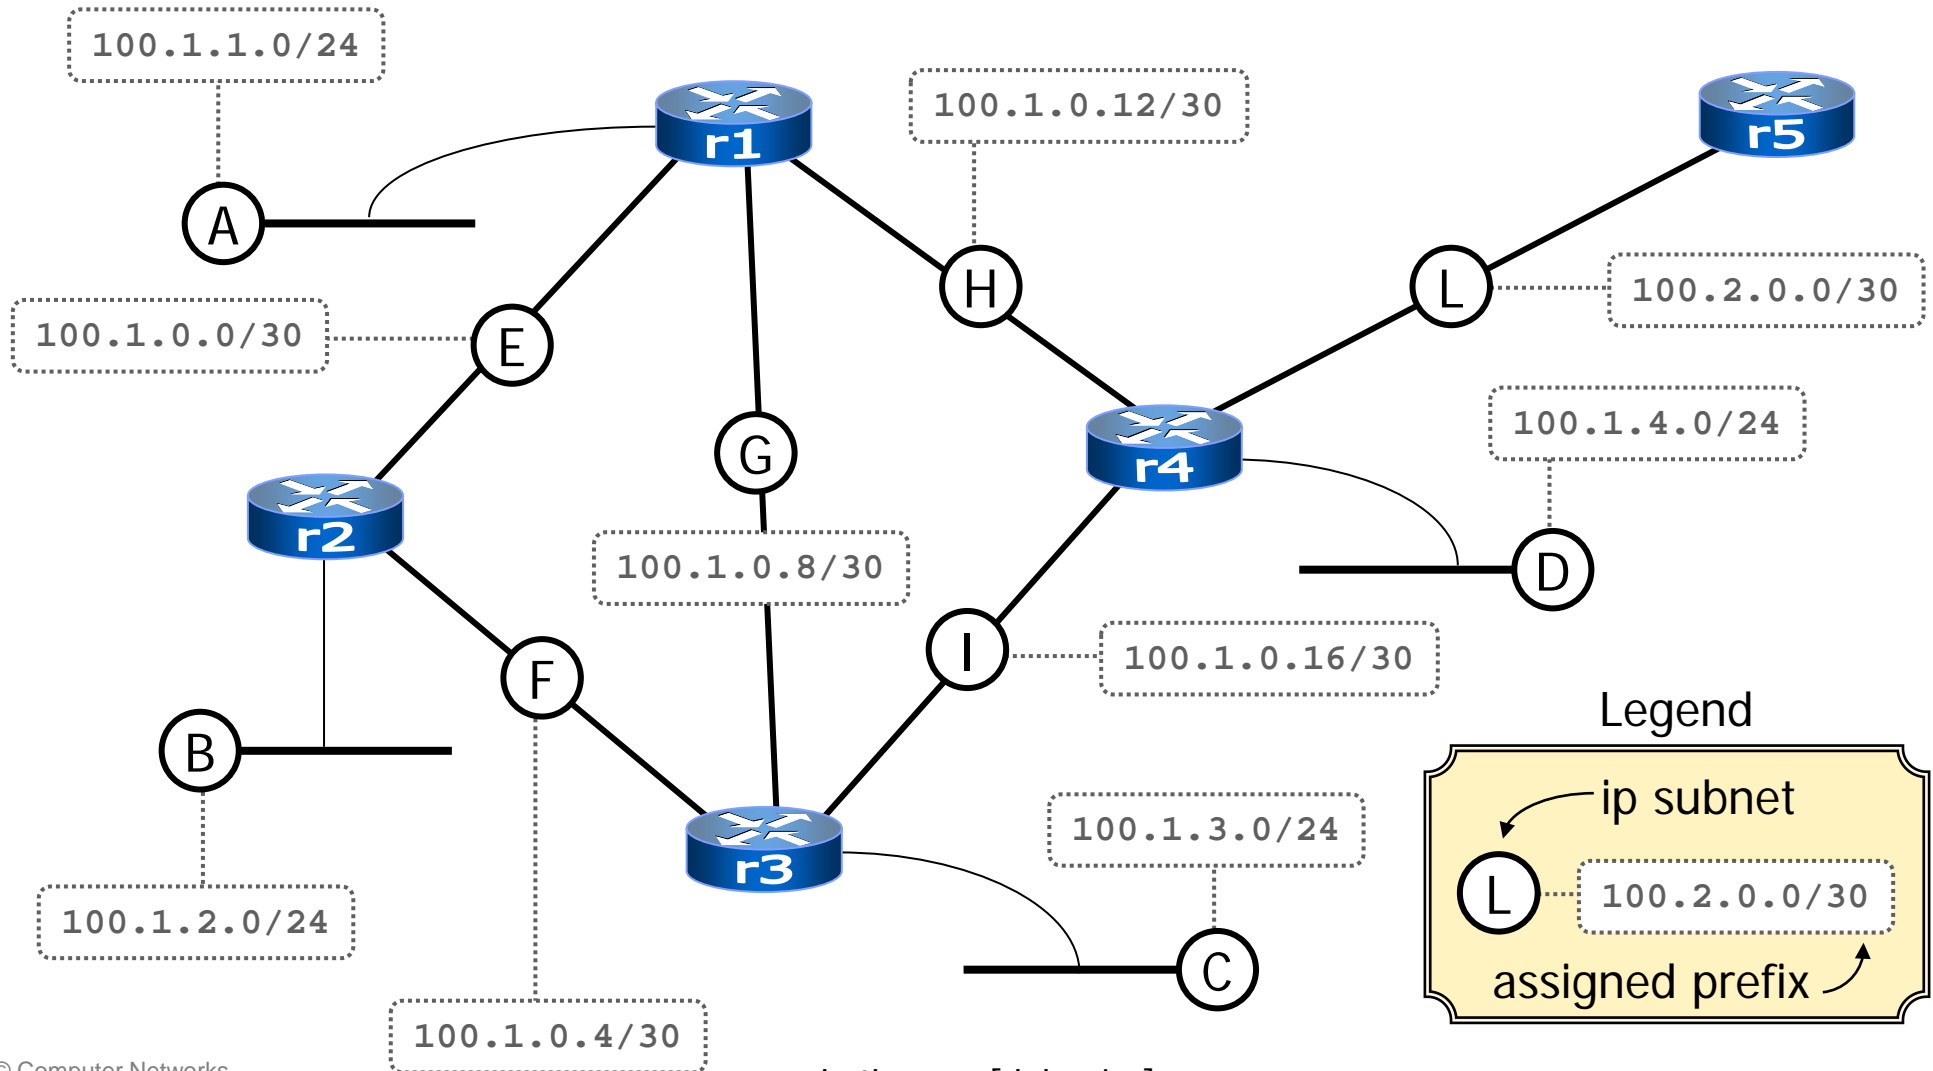
a small network connected to the Internet



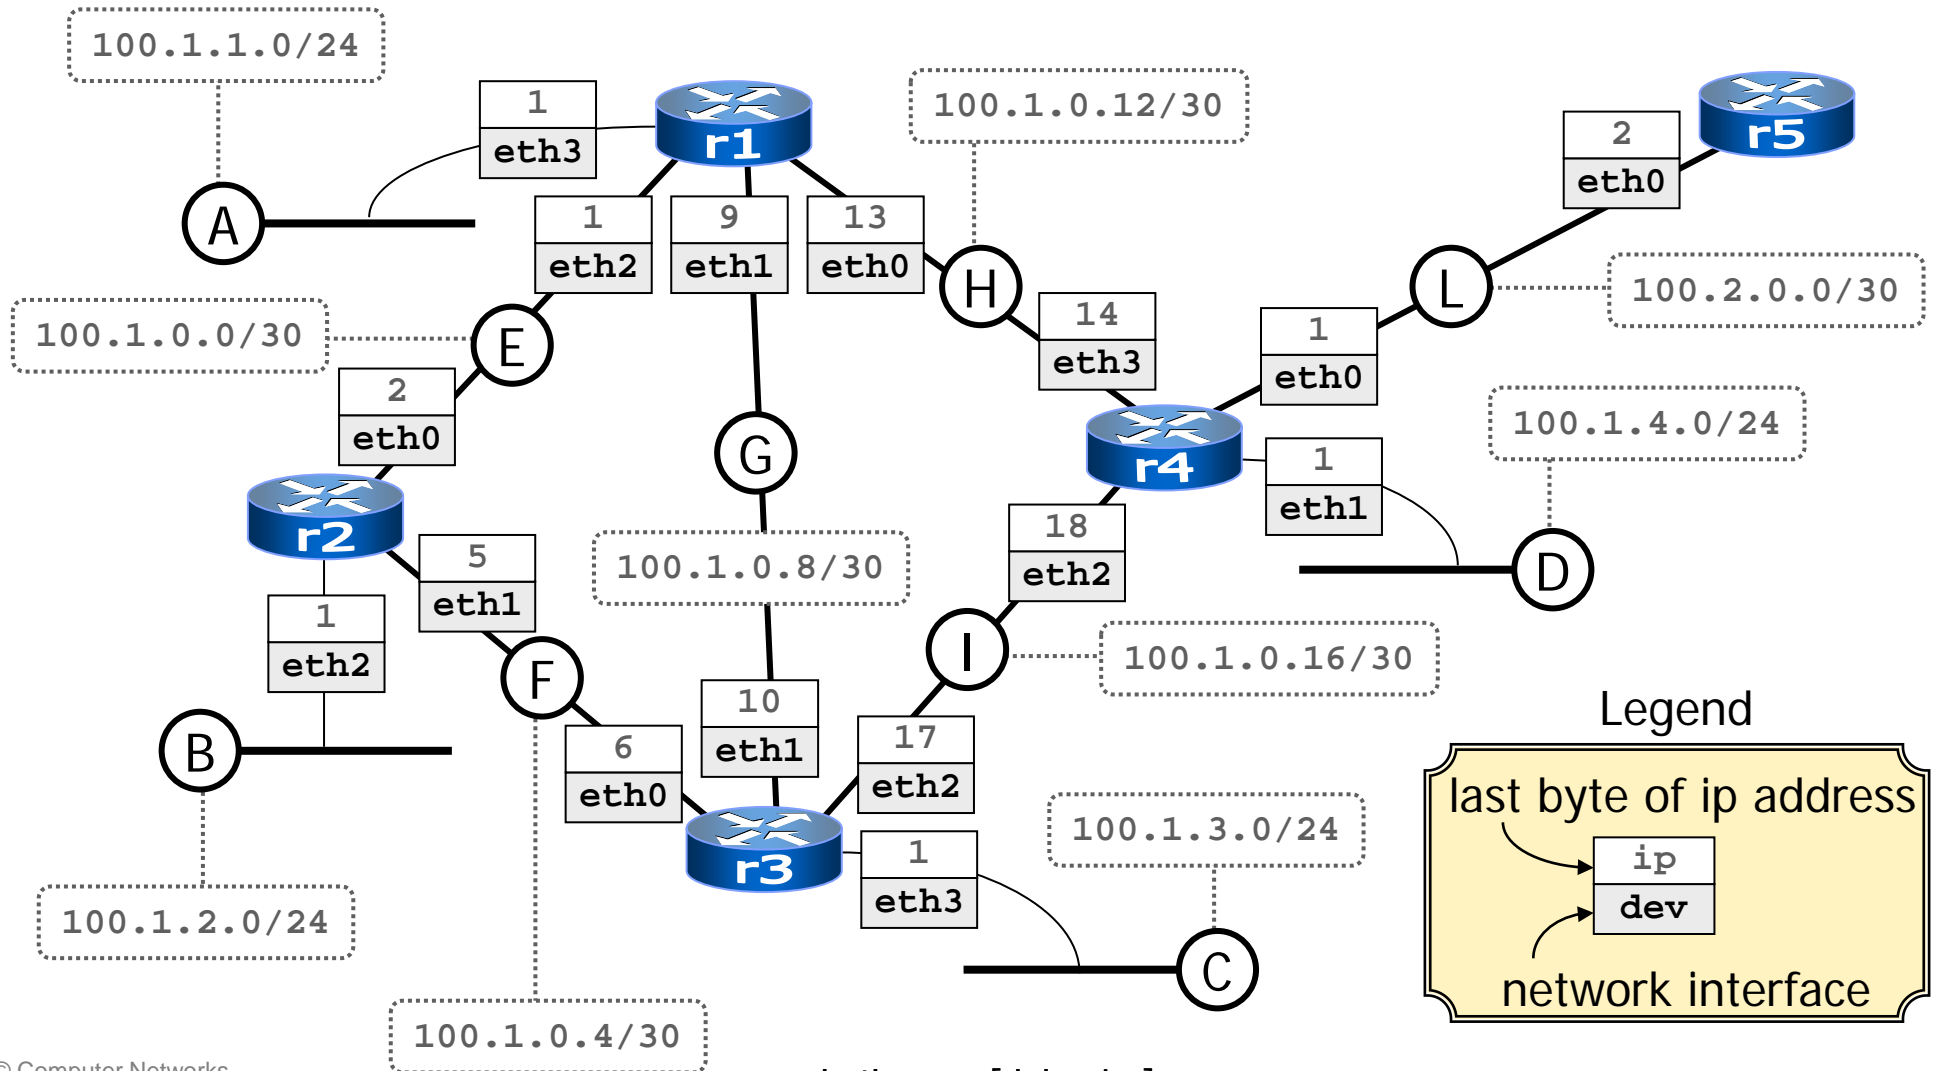
the involved ip subnets



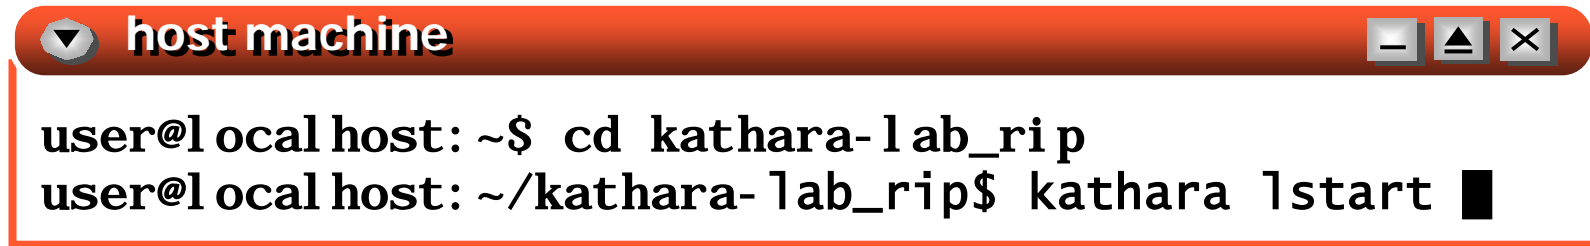
assigning ip numbers to subnets



assigning ip numbers to interfaces



launching the lab script



```
host machine
user@localhost: ~$ cd kathara-lab_rip
user@localhost: ~/kathara-lab_rip$ kathara labstart
```

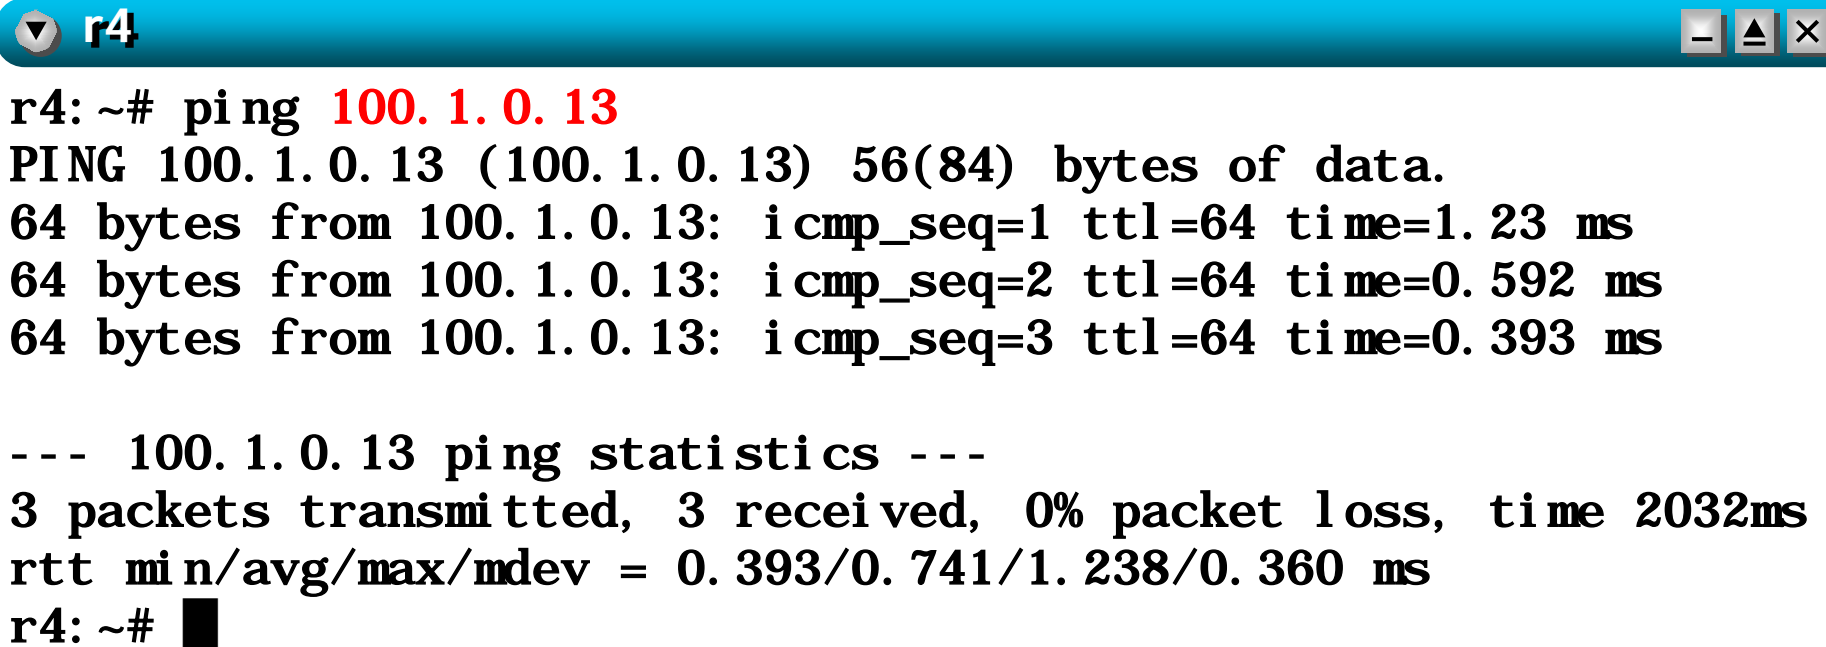
- the lab configuration is such that
 - five virtual hosts are created and connected to the right collision domains (virtual hubs)
 - for each virtual host
 - network interfaces are automatically configured
 - the quagga routing daemon is not automatically started

Lab Scenario Personalization

- Please modify the default scenario in the following way
 - Change the IP address of the `eth0` interface of `r5` to `100.2.0.<LAB-ID>`, where LAB-ID is your personal ID assigned by the lab instructor
- **Note well:** from now-on
 - Command-line commands should reflect the change in addressing, therefore there can be differences in the outputs shown in the manual

checking connectivity

- towards a directly connected destination



```
r4: ~# ping 100.1.0.13
PING 100.1.0.13 (100.1.0.13) 56(84) bytes of data.
64 bytes from 100.1.0.13: icmp_seq=1 ttl=64 time=1.23 ms
64 bytes from 100.1.0.13: icmp_seq=2 ttl=64 time=0.592 ms
64 bytes from 100.1.0.13: icmp_seq=3 ttl=64 time=0.393 ms

--- 100.1.0.13 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2032ms
rtt min/avg/max/mdev = 0.393/0.741/1.238/0.360 ms
r4: ~#
```

checking connectivity

- towards a remote destination

A terminal window with a blue title bar containing a dropdown arrow, the text 'r4', and window control buttons (minimize, maximize, close). The terminal text shows a user at 'r4' running 'ping 100.1.2.1', which results in 'connect: Network is unreachable'. The prompt 'r4: ~#' is followed by a black square cursor.

```
r4: ~# ping 100. 1. 2. 1
connect: Network is unreachable
r4: ~# █
```

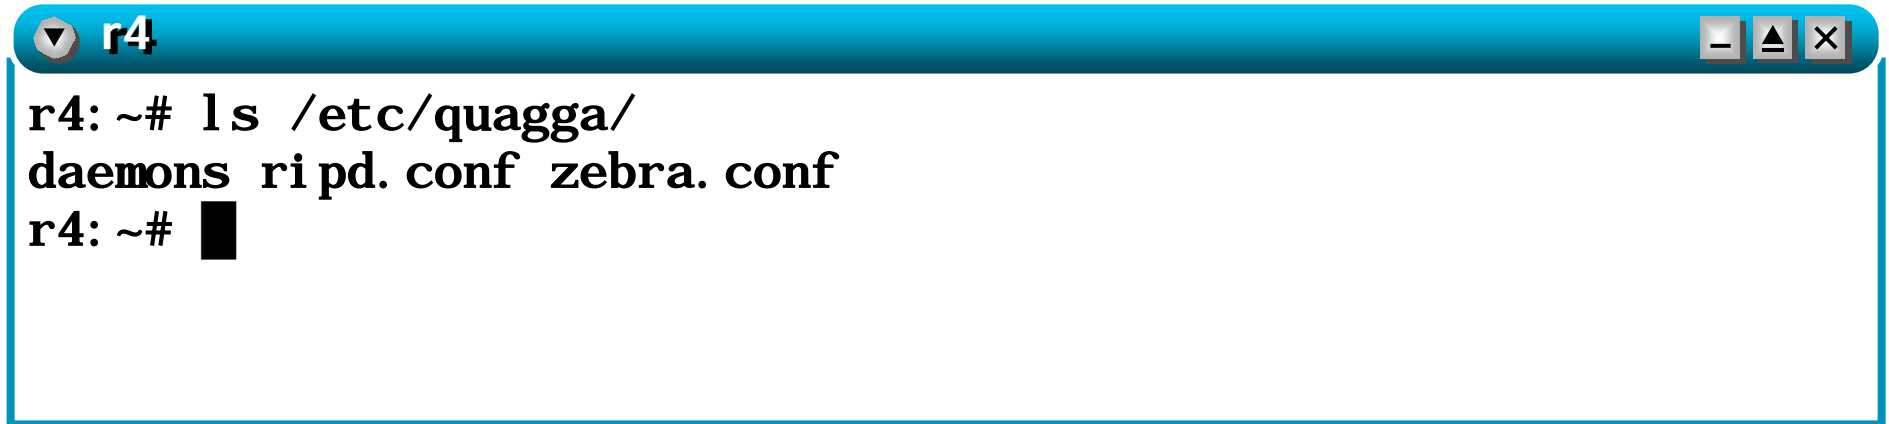
- what's going on?

examining the kernel routing table

```
r4: ~# route
Kernel IP routing table
Destination      Gateway          Genmask          Flags Metric Ref    Use Iface
100.1.0.16       *               255.255.255.252  U        0      0      0 eth2
100.2.0.0         *               255.255.255.252  U        0      0      0 eth0
100.1.0.12        *               255.255.255.252  U        0      0      0 eth3
100.1.4.0         *               255.255.255.0    U        0      0      0 eth1
r4: ~#
```

- since no routing daemon is currently running, only directly connected destinations are known to the router

Inspectig quagga configuration files



```
r4: ~# ls /etc/quagga/  
daemons ripd.conf zebra.conf  
r4: ~#
```

- When quagga is started, each daemon check these files to read the starting configuration

sample zebra configuration file (zebra.conf)

virtual machine

```
pc1:/etc/quagga# less zebra.conf
! *- zebra *-
!
! zebra sample configuration file
!
! $Id: zebra.conf.sample,
Exp $
!
hostname Router
password zebra
enable password zebra
!
! interface lo
zebra.conf
```

the prompt of the zebra interface

the password to connect to the daemon

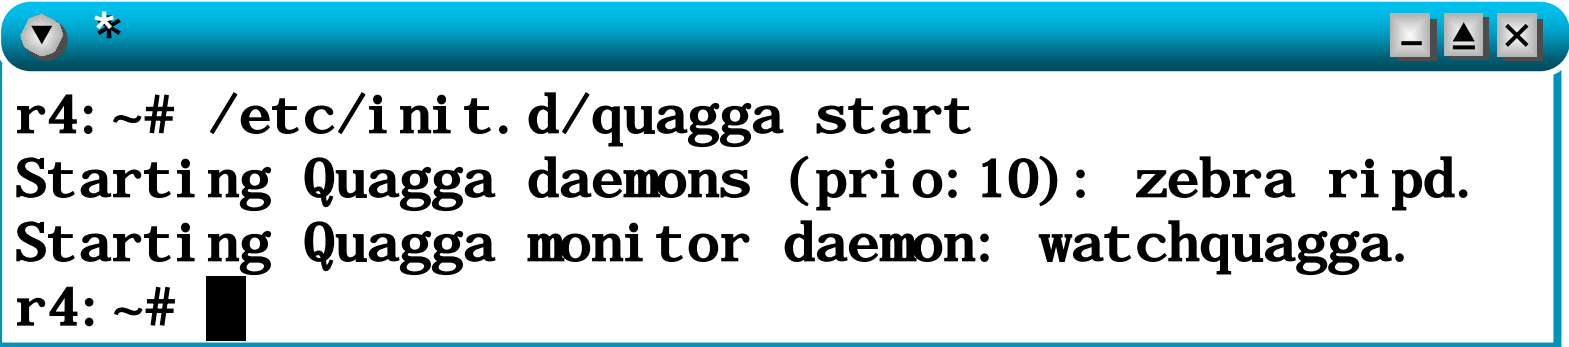
administrative password

daemons file

- Make changes in the daemons files on r1, r2, r3, and r4
 - Location: /etc/quagga/daemons
- Change “**no**” into “**yes**” for **zebra** and **ripd** to start the main and rip daemons in future steps

starting the routing daemons

- on each router (but **r5**) issue the following command:



```
r4: ~# /etc/init.d/quagga start
Starting Quagga daemons (prio: 10): zebra ripd.
Starting Quagga monitor daemon: watchquagga.
r4: ~#
```

Configuring RIP on routers (without router r5)

- telnet to **ripd** and configure rip
- Associate **appropriate** network with RIP process
- Write changes in the **ripd.conf** file

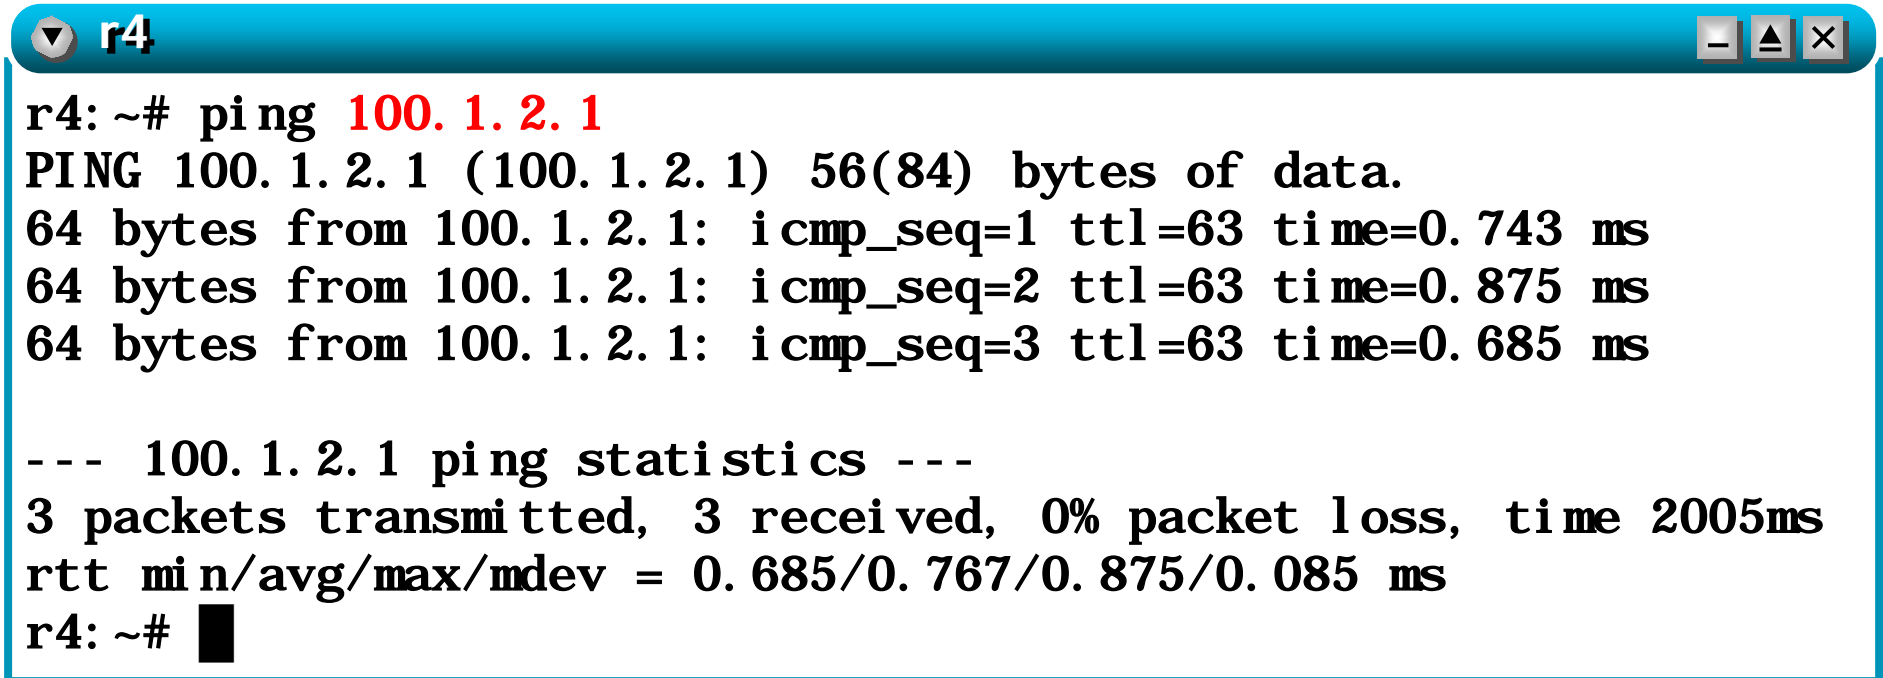
```
r:~# telnet localhost ripd
ripd:> enable
ripd# configure terminal
ripd(config)# router rip
ripd(config-router)# network X.X.X.X/Y
ripd(config-router)# write file
```

Note: Make sure that you have rights to write to quagga configuration files. If not, use the following command:
`chmod a+rw /etc/quagga`

- Be careful when defining interfaces on which RIP runs. Should RIP be spoken on each interface? Control the advertisement of routing information in your network.

checking connectivity (again)

- towards a remote destination



```
r4: ~# ping 100.1.2.1
PING 100.1.2.1 (100.1.2.1) 56(84) bytes of data.
64 bytes from 100.1.2.1: icmp_seq=1 ttl=63 time=0.743 ms
64 bytes from 100.1.2.1: icmp_seq=2 ttl=63 time=0.875 ms
64 bytes from 100.1.2.1: icmp_seq=3 ttl=63 time=0.685 ms

--- 100.1.2.1 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2005ms
rtt min/avg/max/mdev = 0.685/0.767/0.875/0.085 ms
r4: ~# █
```

- after a while, all remote destinations are reachable

checking the routing table

- the routing table is now updated

```
r4: ~# route
Kernel IP routing table
Destination      Gateway          Genmask          Flags Metric Ref    Use Iface
100.1.0.16       *               255.255.255.252  U        0      0      0 eth2
100.1.0.0        100.1.0.13      255.255.255.252  UG       2      0      0 eth3
100.1.0.4        100.1.0.17      255.255.255.252  UG       2      0      0 eth2
100.2.0.0        *               255.255.255.252  U        0      0      0 eth0
100.1.0.8        100.1.0.17      255.255.255.252  UG       2      0      0 eth2
100.1.0.12       *               255.255.255.252  U        0      0      0 eth3
100.1.4.0        *               255.255.255.0    U        0      0      0 eth1
100.1.2.0        100.1.0.17      255.255.255.0    UG       3      0      0 eth2
100.1.3.0        100.1.0.17      255.255.255.0    UG       2      0      0 eth2
100.1.1.0        100.1.0.13      255.255.255.0    UG       2      0      0 eth3
r4: ~# █
```

a look at ripv2 packets

- let's sniff ripv2 packets

```
r4: ~# tcpdump -i eth2 -v -n -s 1518
```

display packet details
(enable full protocol decoding)

don't resolve numbers
to names

sniff entire ethernet
packets (by default, only
the first 68 bytes are
captured)

a look at ripv2 packets

- let's sniff ripv2 packets

```
r4: ~# tcpdump -i eth2 -v -n -s 1518
tcpdump: listening on eth2, link-type EN10MB (Ethernet), capture size 1518
bytes
16:47:48.333986 IP (tos 0x0, ttl 1, id 0, offset 0, flags [DF], length: 152)
100.1.0.17.520 > 224.0.0.9.520: [udp sum ok]
    RIPv2, Response, length: 124, routes: 6
        AFI: IPv4:      100.1.0.0/30, tag 0x0000, metric: 2, next-hop: self
        AFI: IPv4:      100.1.0.4/30, tag 0x0000, metric: 1, next-hop: self
        AFI: IPv4:      100.1.0.8/30, tag 0x0000, metric: 1, next-hop: self
        AFI: IPv4:      100.1.1.0/24, tag 0x0000, metric: 2, next-hop: self
        AFI: IPv4:      100.1.2.0/24, tag 0x0000, metric: 2, next-hop: self
        AFI: IPv4:      100.1.3.0/24, tag 0x0000, metric: 1, next-hop: self

1 packets captured
1 packets received by filter
0 packets dropped by kernel
r4: ~#
```


a traceroute

r4

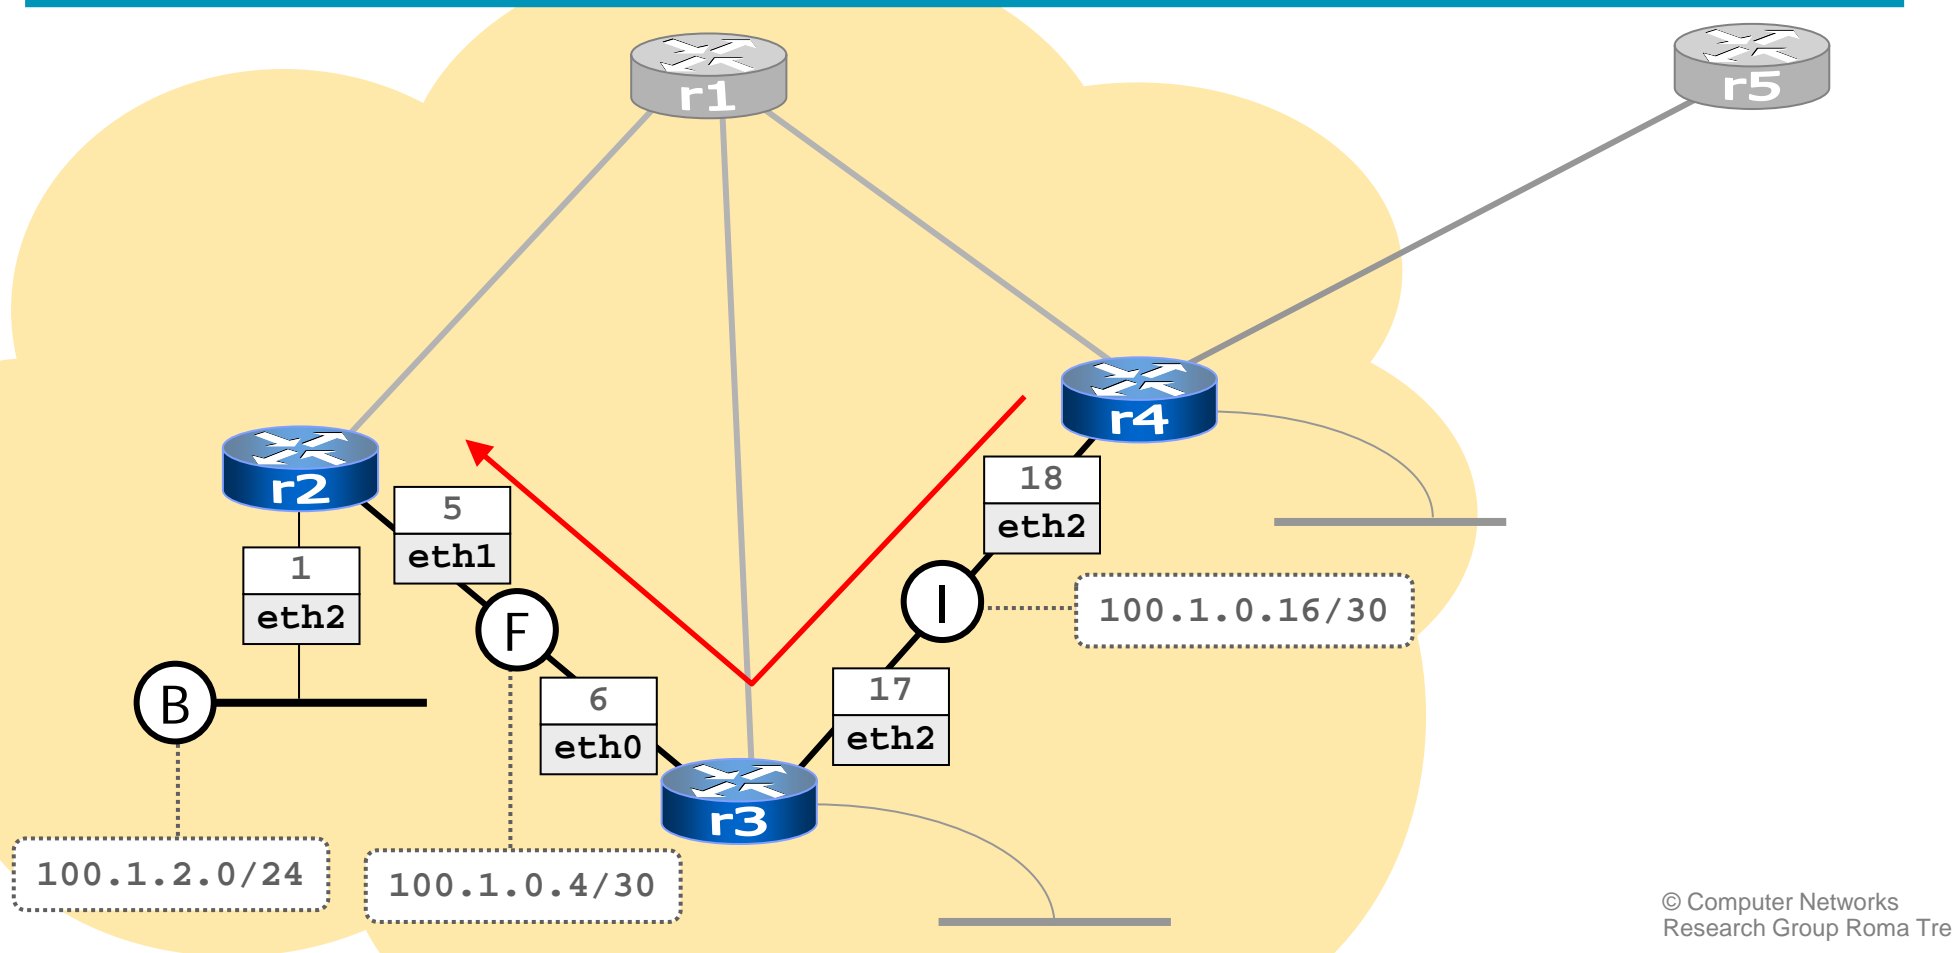
```
r4: ~# traceroute 100.1.2.1
```

```
traceroute to 100.1.2.1 (100.1.2.1), 64 hops max, 40 byte packets
```

```
1 100.1.0.17 (100.1.0.17) 10 ms 3 ms 1 ms
```

```
2 100.1.2.1 (100.1.2.1) 15 ms 1 ms 1 ms
```

```
r4: ~#
```



inspecting the rip routing table

r4

```
r4: ~# telnet localhost ripd
```

```
.....
```

```
Password: zebra
```

```
ripd> show ip rip
```

```
Codes: R - RIP, C - connected, O - OSPF, B - BGP
```

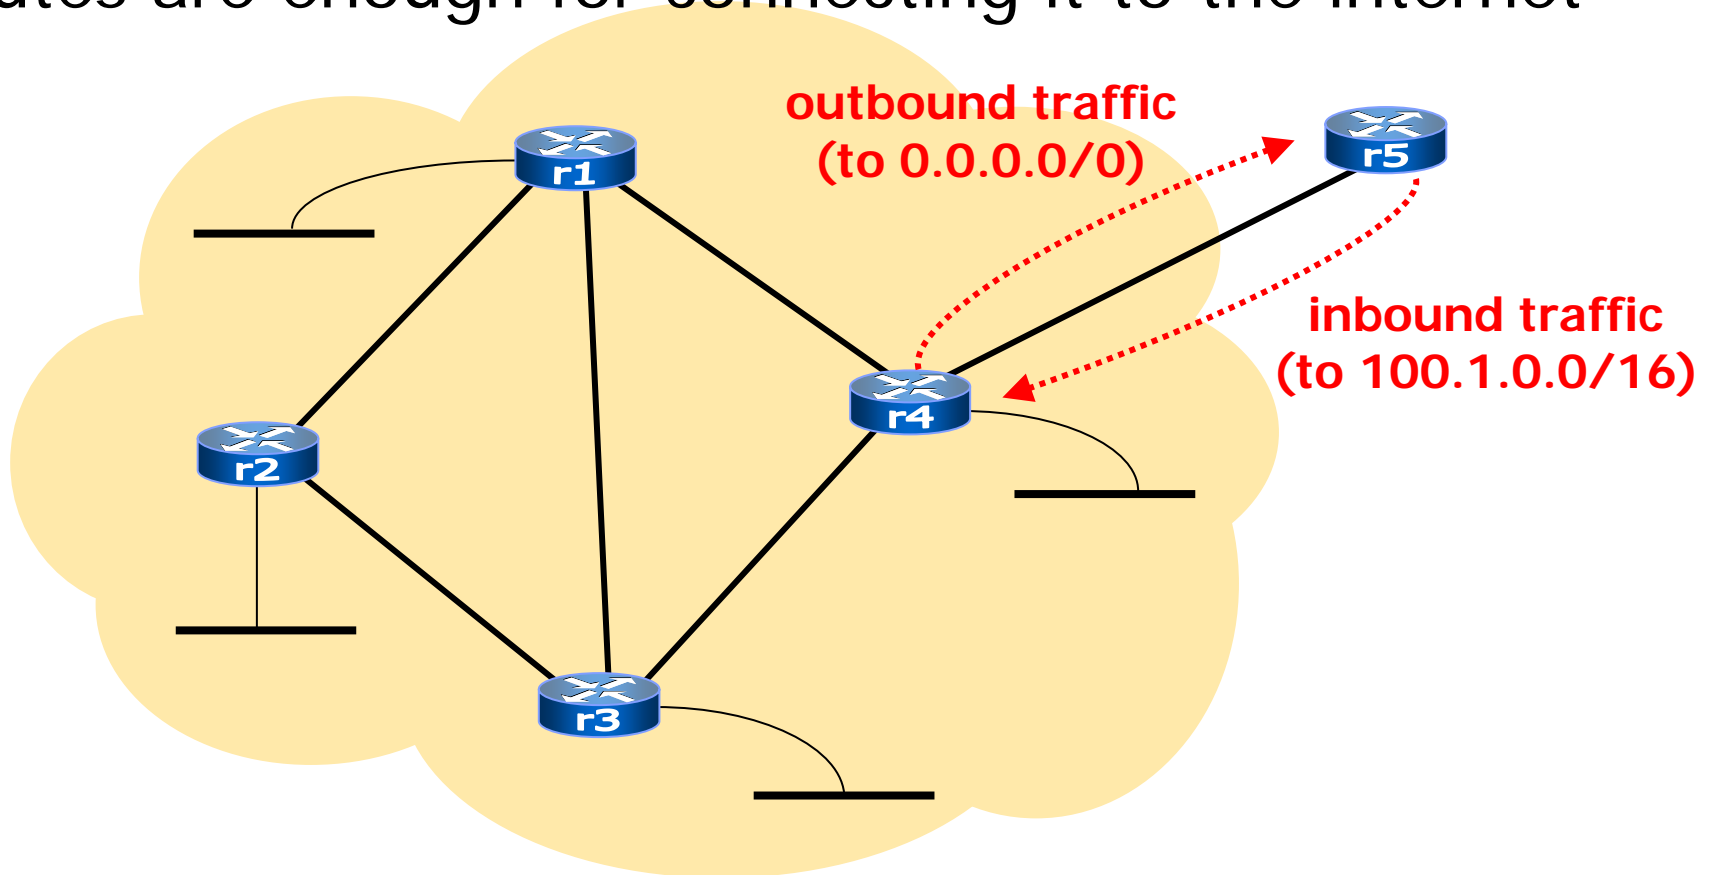
```
(n) - normal, (s) - static, (d) - default, (r) - redistribute,  
(i) - interface
```

	Network	Next Hop	Metric	From	Time
R(n)	100. 1. 0. 0/30	100. 1. 0. 13	2	100. 1. 0. 13	02: 43
R(n)	100. 1. 0. 4/30	100. 1. 0. 17	2	100. 1. 0. 17	02: 46
R(n)	100. 1. 0. 8/30	100. 1. 0. 17	2	100. 1. 0. 17	02: 46
C(i)	100. 1. 0. 12/30	0. 0. 0. 0	1	self	
C(i)	100. 1. 0. 16/30	0. 0. 0. 0	1	self	
R(n)	100. 1. 1. 0/24	100. 1. 0. 13	2	100. 1. 0. 13	02: 43
R(n)	100. 1. 2. 0/24	100. 1. 0. 17	3	100. 1. 0. 17	02: 46
R(n)	100. 1. 3. 0/24	100. 1. 0. 17	2	100. 1. 0. 17	02: 46
C(i)	100. 1. 4. 0/24	0. 0. 0. 0	1	self	
C(r)	100. 2. 0. 0/30	0. 0. 0. 0	1	self	

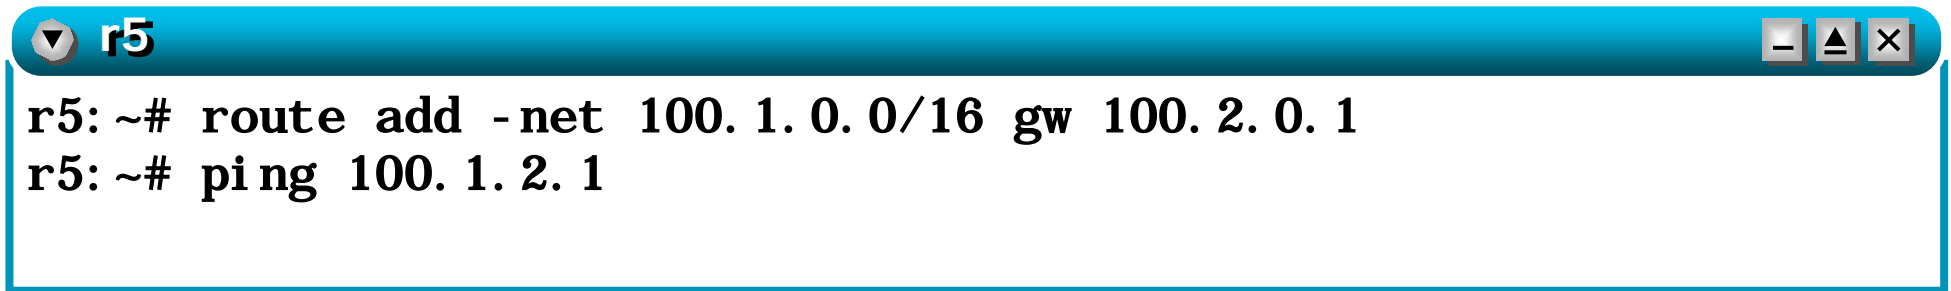
```
ripd> █
```

static routing

- our network is a **stub network** (i.e., it has just one connection to an external router, **r5**); hence, static routes are enough for connecting it to the internet



adding a static route to r5

A terminal window with a blue title bar containing a dropdown arrow and the text 'r5'. On the right side of the title bar are three window control icons: minimize, maximize, and close. The terminal area has a white background and contains two lines of text: 'r5: ~# route add -net 100. 1. 0. 0/16 gw 100. 2. 0. 1' and 'r5: ~# ping 100. 1. 2. 1'.

```
r5: ~# route add -net 100. 1. 0. 0/16 gw 100. 2. 0. 1
r5: ~# ping 100. 1. 2. 1
```

- Let **eth0** of **r4** to participate in RIP
- Add an appropriate configuration to ripd
- Hints:
 - ripd(config-router)# redistribute connected
 - ripd(config-router)# write file

checking connectivity

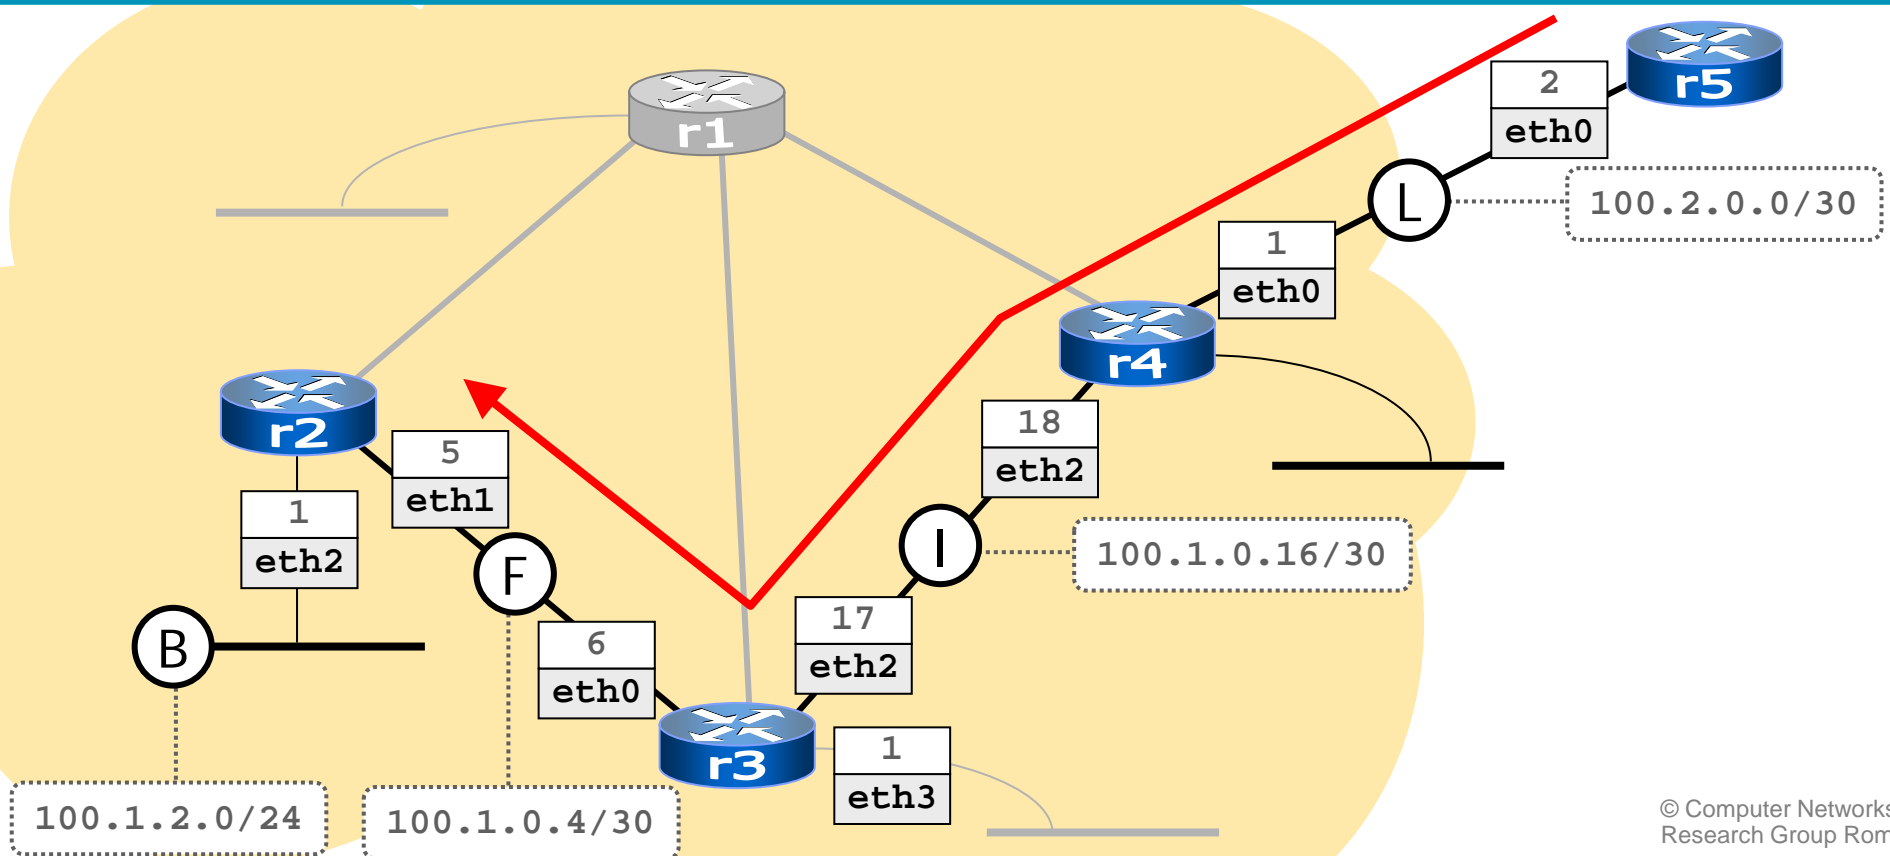
r5

```
r5: ~# traceroute 100.1.2.1
```

```
traceroute to 100.1.2.1 (100.1.2.1), 64 hops max, 40 byte packets
```

```
1  100.2.0.1 (100.2.0.1)  75 ms  1 ms  2 ms
2  100.1.0.17 (100.1.0.17) 7 ms   1 ms  1 ms
3  100.1.2.1 (100.1.2.1) 24 ms  3 ms  1 ms
```

```
r5: ~# █
```



Reporting

- Please deliver the following items to the UPEL system using your account
 1. A photocopy or a screenshot showing the output of the **traceroute 100.1.2.1** command executed on r5 (from the previous slide)

configuring r4

■ step 1: configuring the default route

r4

```
r4: ~# route add default gw 100.2.0.2
```

```
r4: ~# route
```

```
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
100.1.0.16	*	255.255.255.252	U	0	0	0	eth2
100.1.0.0	100.1.0.13	255.255.255.252	UG	2	0	0	eth3
100.1.0.4	100.1.0.17	255.255.255.252	UG	2	0	0	eth2
100.2.0.0	*	255.255.255.252	U	0	0	0	eth0
100.1.0.8	100.1.0.17	255.255.255.252	UG	2	0	0	eth2
100.1.0.12	*	255.255.255.252	U	0	0	0	eth3
100.1.4.0	*	255.255.255.0	U	0	0	0	eth1
100.1.2.0	100.1.0.17	255.255.255.0	UG	3	0	0	eth2
100.1.3.0	100.1.0.17	255.255.255.0	UG	2	0	0	eth2
100.1.1.0	100.1.0.13	255.255.255.0	UG	2	0	0	eth3
default	100.2.0.2	0.0.0.0	UG	0	0	0	eth0

```
r4: ~#
```

This configuration adds a default gw for r4

configuring r4

- step 2: propagating the default route into rip

r4

```
r4: ~# telnet localhost ripd
Trying 127.0.0.1...
Connected to r4.
Escape character is '^['.
```

```
Hello, this is Quagga (version 1.0.20160315).
Copyright 1996-2005 Kunihiro Ishiguro, et al.
```

User Access Verification

Password: zebra

```
ripd> enable
```

Password: zebra

```
ripd# configure terminal
```

```
ripd(config)# router rip
```

```
ripd(config-router)# route 0.0.0.0/0
```

```
ripd(config-router)# quit
```

```
ripd(config)# quit
```

```
ripd# disable
```

```
ripd> exit
```

This configuration propagates a default gw for other rip speakers than r4

work as a privileged user

begin configuration

configure the rip protocol

statically configure the default route

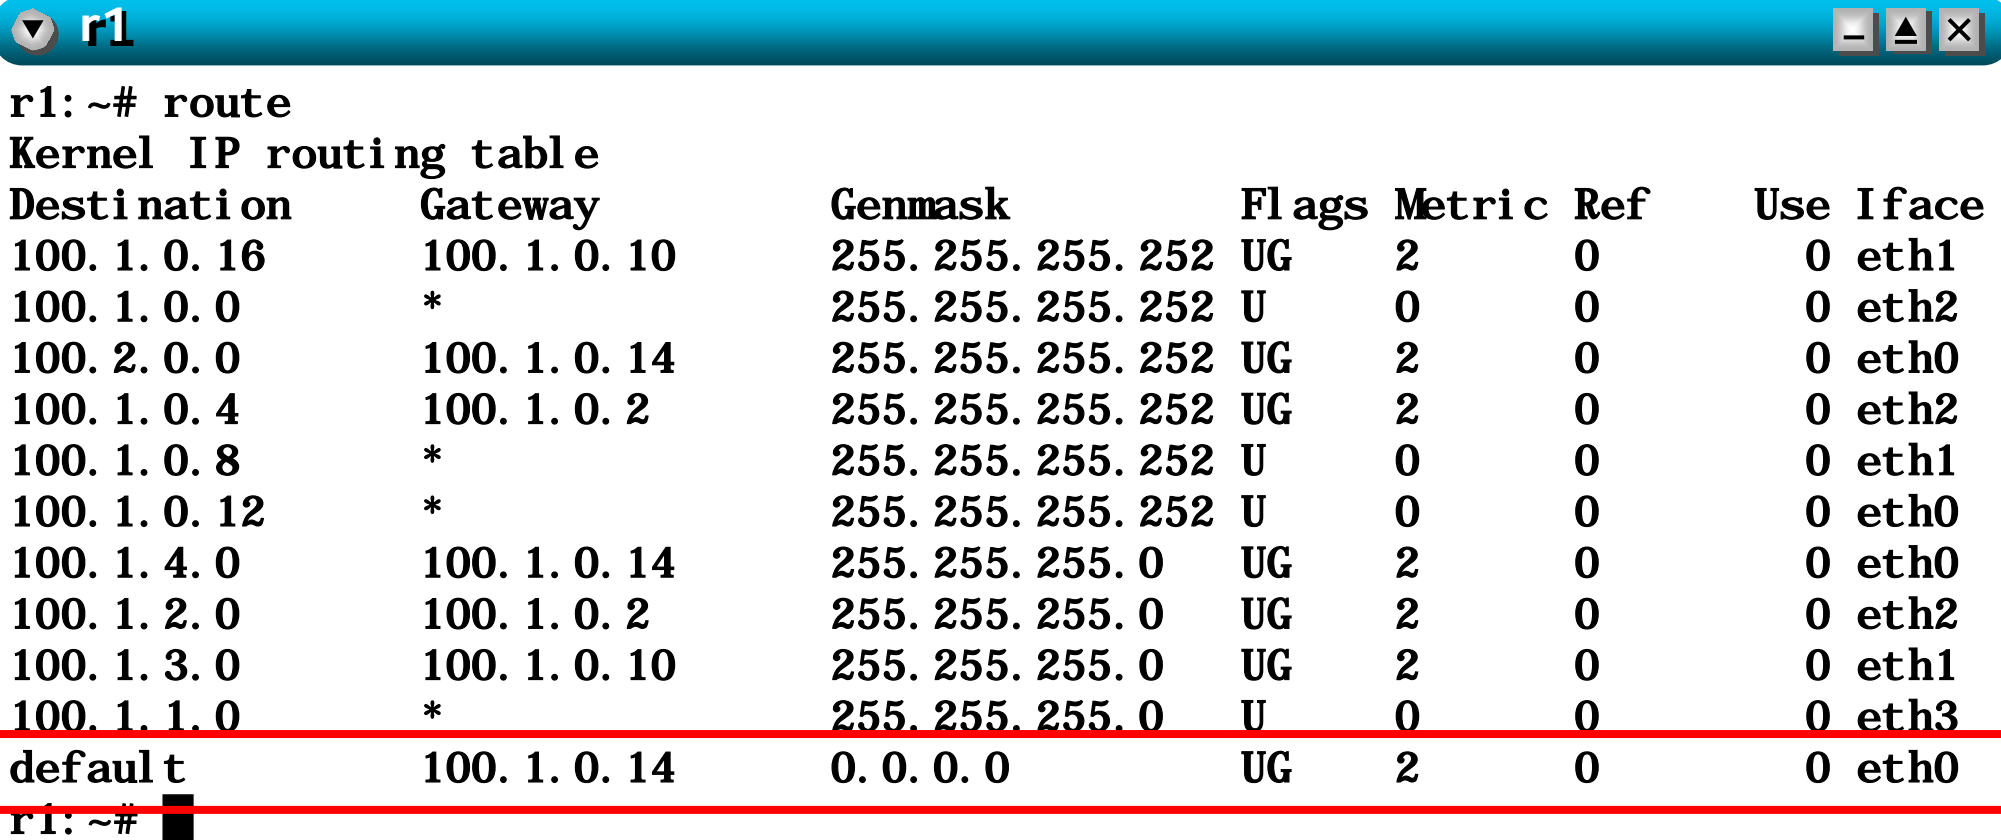
end of rip configuration

end configuration

abandon privileges

the default route

- after a while, the default route has been injected (via rip) into the network



```
r1: ~# route
Kernel IP routing table
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	Iface
100.1.0.16	100.1.0.10	255.255.255.252	UG	2	0	0	eth1
100.1.0.0	*	255.255.255.252	U	0	0	0	eth2
100.2.0.0	100.1.0.14	255.255.255.252	UG	2	0	0	eth0
100.1.0.4	100.1.0.2	255.255.255.252	UG	2	0	0	eth2
100.1.0.8	*	255.255.255.252	U	0	0	0	eth1
100.1.0.12	*	255.255.255.252	U	0	0	0	eth0
100.1.4.0	100.1.0.14	255.255.255.0	UG	2	0	0	eth0
100.1.2.0	100.1.0.2	255.255.255.0	UG	2	0	0	eth2
100.1.3.0	100.1.0.10	255.255.255.0	UG	2	0	0	eth1
100.1.1.0	*	255.255.255.0	U	0	0	0	eth3
default	100.1.0.14	0.0.0.0	UG	2	0	0	eth0

```
r1: ~#
```

checking connectivity

▼ r1


any (even non-existing) destination

```
r1: ~# ping 193.204.161.1
PING 193.204.161.1 (193.204.161.1) 56(84) bytes of data.
From 100.2.0.2 icmp_seq=1 Destination Net Unreachable
From 100.2.0.2 icmp_seq=2 Destination Net Unreachable
```

```
--- 193.204.161.1 ping statistics ---
```

```
2 packets transmitted, 0 received, +2 errors, 100% packet loss,
time 999ms
```

```
r1: ~#
```



checking connectivity

- **r5** is actually receiving echo request packets

```
r5: ~# tcpdump -i eth0 -n -s 1518
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on eth0, link-type EN10MB (Ethernet), capture size 1518 bytes
11:38:43.822503 arp who-has 100.2.0.2 tell 100.2.0.1
11:38:43.824221 arp reply 100.2.0.2 is-at fe:fd:64:02:00:02
11:38:43.825890 IP 100.1.0.13 > 193.204.161.1: icmp 64: echo request seq 1
11:38:43.827139 IP 100.2.0.2 > 100.1.0.13: icmp 92: net 193.204.161.1
unreachable
11:38:44.841566 IP 100.1.0.13 > 193.204.161.1: icmp 64: echo request seq 2
11:38:44.841651 IP 100.2.0.2 > 100.1.0.13: icmp 92: net 193.204.161.1
unreachable

6 packets captured
6 packets received by filter
0 packets dropped by kernel
r5: ~#
```


Reporting

- Please deliver the following items to the UPEL system using your account
 1. A photocopy or a screenshot showing the output of the **tcpdump -i eth0 -n -s 1518** command executed on r5 (from the previous slide)

shutting down an interface

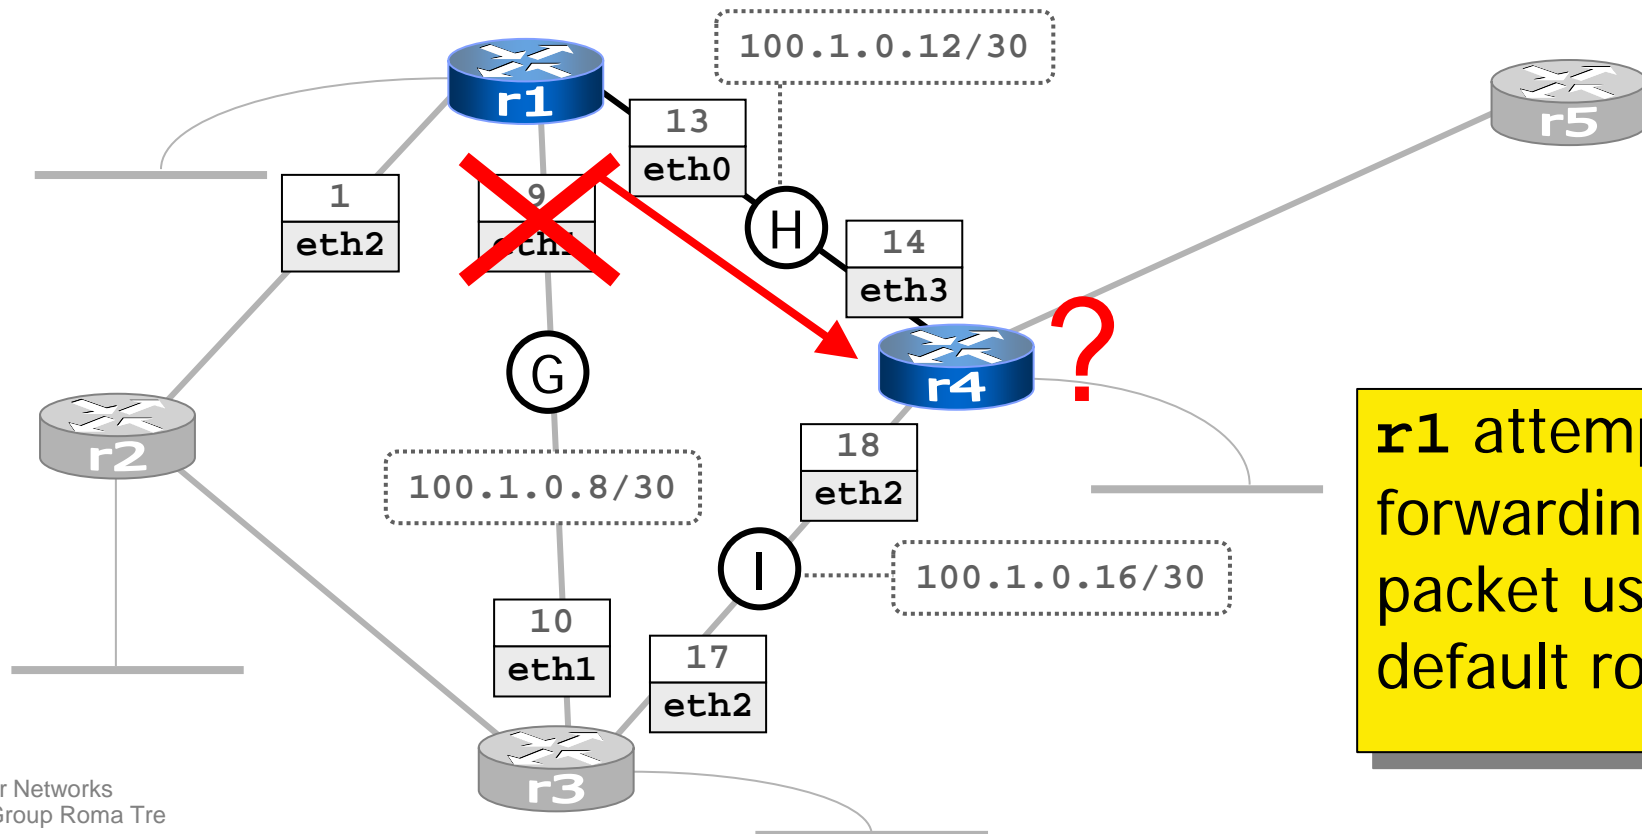
▼ r1

```
r1:~# traceroute 100.1.0.10
traceroute to 100.1.0.10 (100.1.0.10), 64 hops max, 40 byte packets
 1  100.1.0.10 (100.1.0.10)  24 ms  1 ms  1 ms
r1:~# ifconfig eth1 down
```



shutting down an interface

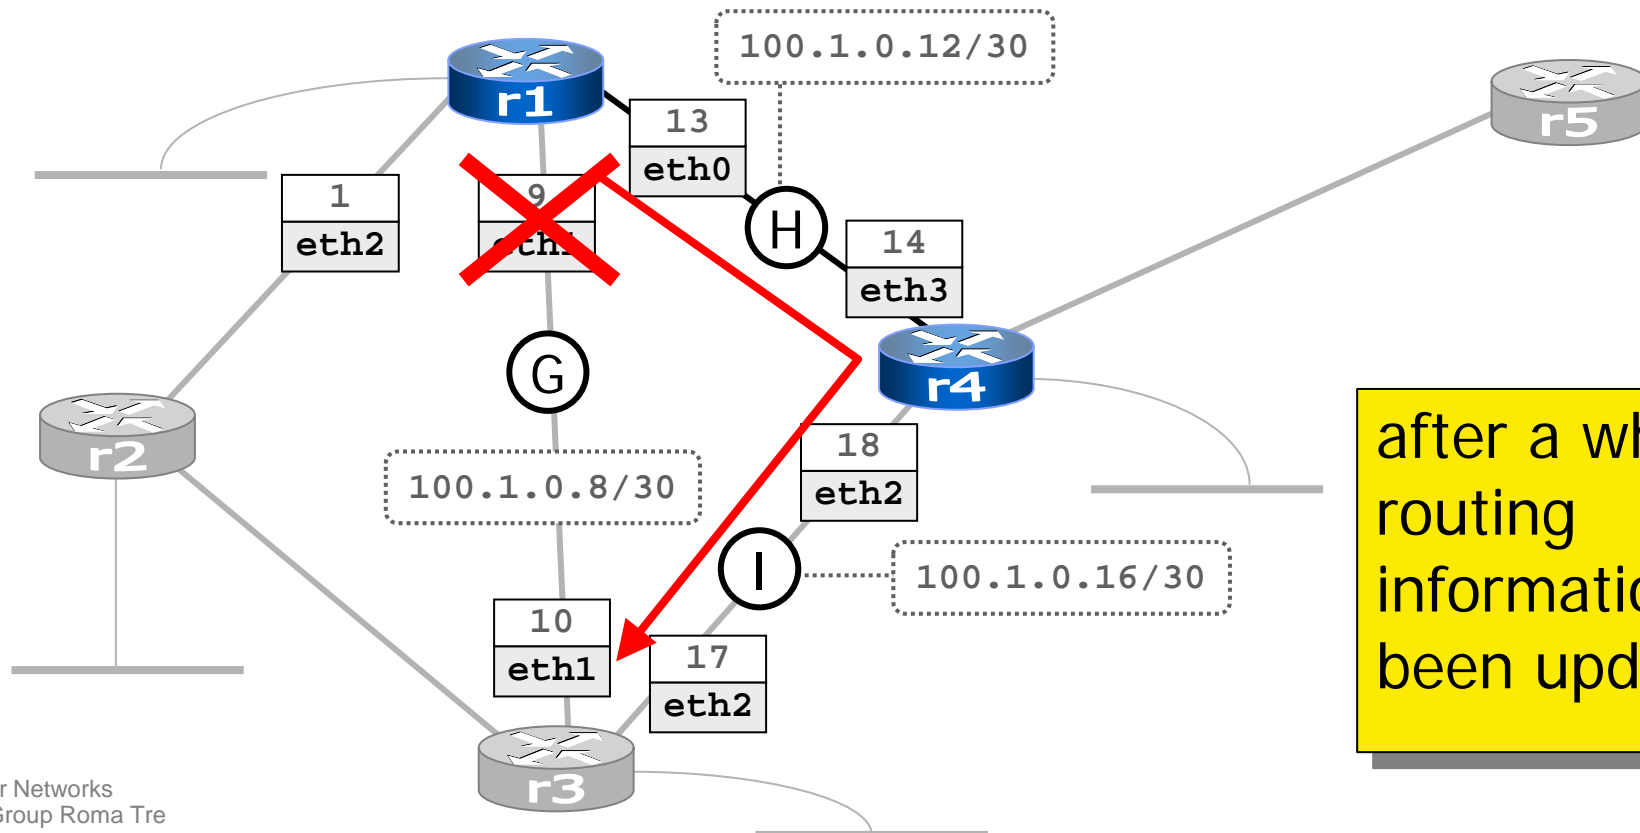
```
r1:~# traceroute 100.1.0.10
traceroute to 100.1.0.10 (100.1.0.10), 64 hops max, 40 byte packets
 1  100.1.0.14 (100.1.0.14)  1 ms  1 ms  1 ms
 2  * * *
 3  * * * █
```



r1 attempts forwarding the packet using the default route

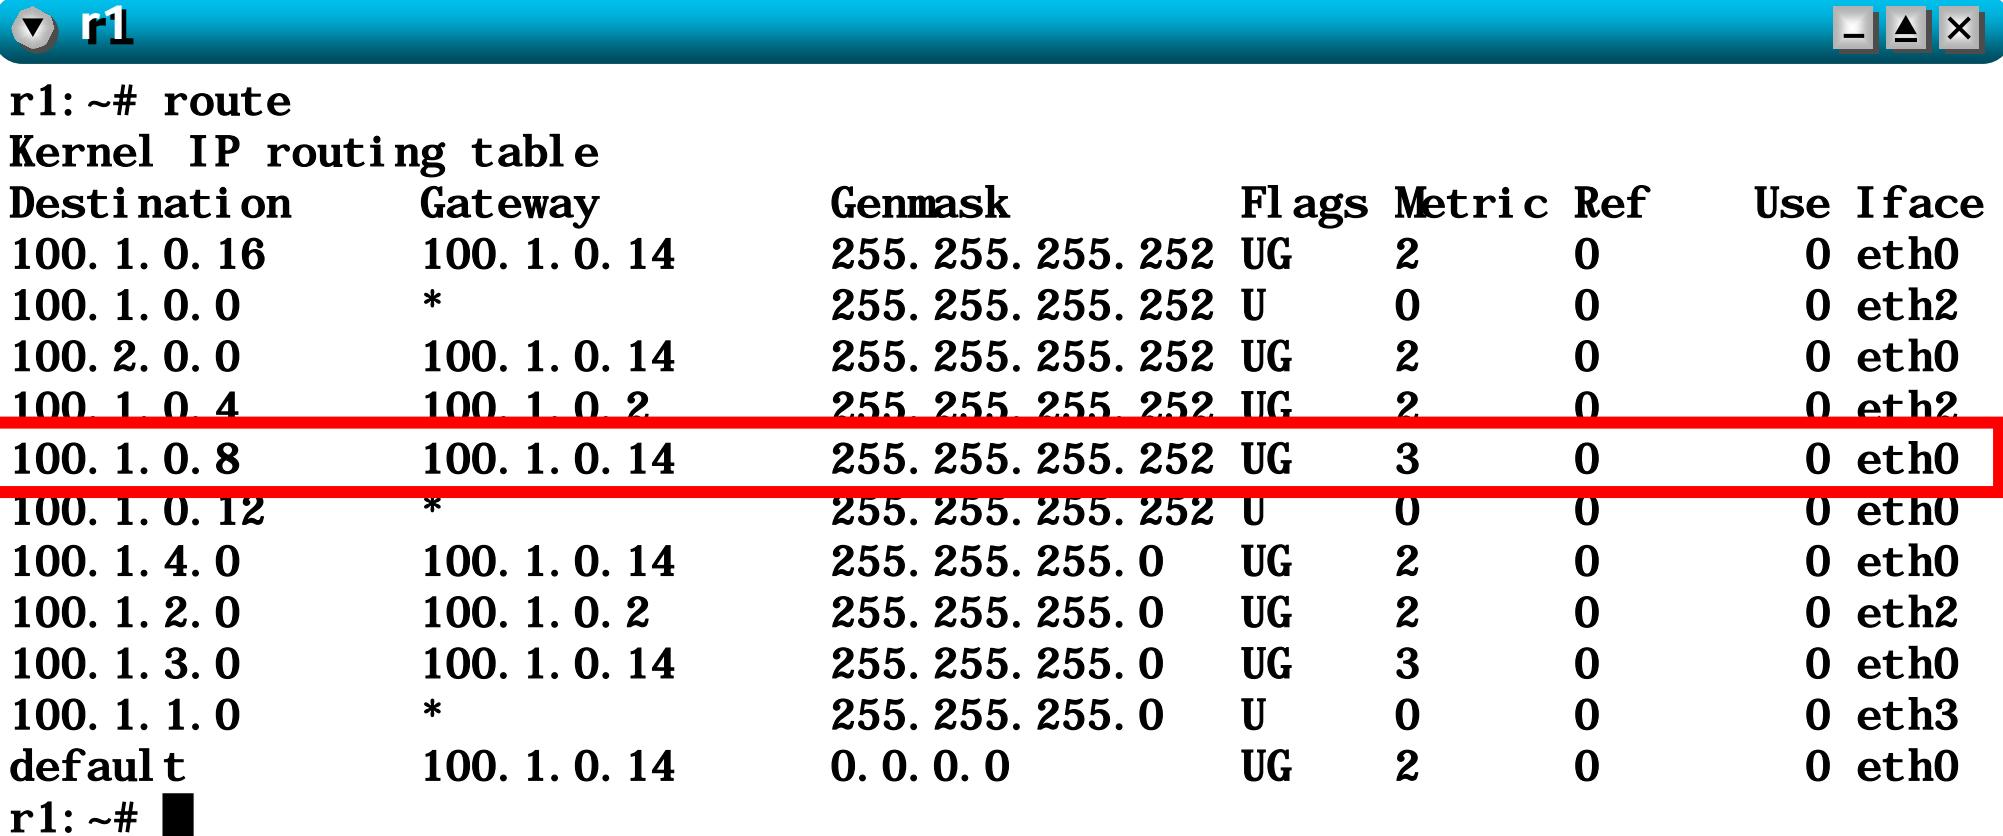
shutting down an interface

```
r1: ~# traceroute 100.1.0.10
traceroute to 100.1.0.10 (100.1.0.10), 64 hops max, 40 byte packets
 1  100.1.0.14 (100.1.0.14)  1 ms  1 ms  1 ms
 2  100.1.0.10 (100.1.0.10)  5 ms  2 ms  1 ms
r1: ~# █
```



shutting down an interface

- **r1**'s routing table has been updated



```
r1: ~# route
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use I face
100.1.0.16       100.1.0.14      255.255.255.252 UG      2      0      0 eth0
100.1.0.0         *               255.255.255.252 U        0      0      0 eth2
100.2.0.0         100.1.0.14      255.255.255.252 UG      2      0      0 eth0
100.1.0.4         100.1.0.2       255.255.255.252 UG      2      0      0 eth2
100.1.0.8         100.1.0.14      255.255.255.252 UG      3      0      0 eth0
100.1.0.12        *               255.255.255.252 U        0      0      0 eth0
100.1.4.0         100.1.0.14      255.255.255.0   UG      2      0      0 eth0
100.1.2.0         100.1.0.2       255.255.255.0   UG      2      0      0 eth2
100.1.3.0         100.1.0.14      255.255.255.0   UG      3      0      0 eth0
100.1.1.0         *               255.255.255.0   U        0      0      0 eth3
default          100.1.0.14      0.0.0.0         UG      2      0      0 eth0
r1: ~#
```

Destination	Gateway	Genmask	Flags	Metric	Ref	Use	I face
100.1.0.16	100.1.0.14	255.255.255.252	UG	2	0	0	eth0
100.1.0.0	*	255.255.255.252	U	0	0	0	eth2
100.2.0.0	100.1.0.14	255.255.255.252	UG	2	0	0	eth0
100.1.0.4	100.1.0.2	255.255.255.252	UG	2	0	0	eth2
100.1.0.8	100.1.0.14	255.255.255.252	UG	3	0	0	eth0
100.1.0.12	*	255.255.255.252	U	0	0	0	eth0
100.1.4.0	100.1.0.14	255.255.255.0	UG	2	0	0	eth0
100.1.2.0	100.1.0.2	255.255.255.0	UG	2	0	0	eth2
100.1.3.0	100.1.0.14	255.255.255.0	UG	3	0	0	eth0
100.1.1.0	*	255.255.255.0	U	0	0	0	eth3
default	100.1.0.14	0.0.0.0	UG	2	0	0	eth0

about **redistribute connected**

- by default (i.e., without further configuration) RIP already propagates information about directly connected subnets...
...attached to RIP-speaking interfaces only
- **redistribute connected** forces RIP to propagate information about *all* connected subnets
- the semantic of **redistribute connected** applies to all routing protocols
- the default behavior does not: some protocols (e.g., bgp) are lazier, and do not propagate anything unless explicitly told to do so