

BGP: Local Preference attribute

BGP: Atrybut lokalnej preferencji

- Celem zajęć jest ugruntowanie wiedzy dotyczącej atrybutu **Local Preference**, pozwalającego na zarządzanie ruchem w kierunku wychodzącym z danego systemu AS.
- W ramach zajęć zostanie uruchomione 6 ruterów z podstawową konfiguracją sieciową oraz częściową konfiguracją BGP.
- W ramach zajęć należy:
 - Sprawdzić osiągalność sieci i przebieg tras w podstawowej konfiguracji BGP (bez ustawionego atrybutu **LocPref**).
 - Skonfigurować właściwe wartości atrybutu **LocPref** na odpowiednich ruterach.
 - Sprawdzić osiągalność sieci i przebieg tras po zastosowaniu atrybutu **LocPref**.

Traffic Engineering with BGP

BGP: Outbound Traffic Control

Authors: Piotr Pacyna, Katarzyna Kosek-Szott

Overview

Objective: show propagation of prefixes between ISPs in BGP multi-homing scenario

- ISP can
 - manipulate prefix during propagation/processing to satisfy its own goals regarding path selection
 - implement its own policies for inbound and outbound flows
- Specifically, we will see that
 - **Part I (LocalPref):** ISP **can**
 - influence path selection process carried out by their own BGP routers
 - influence the path for **outbound traffic**

Outline

- General idea: insert additional information into **BGP Update** messages
 - Extra information propagated with the prefix and evaluated by BGP routers during the BGP decision process
 - Appropriate routes installed in routing tables of BGP routers
 - Inter-domain traffic routed following the chosen routes

Part I

Outbound traffic control with LocPref

Part I Outline

■ Goal

- Influence on path selection for **outbound traffic**, i.e., traffic originated in a local AS, destined to a network in another AS

■ Idea

- Path selection for outbound traffic carried out in a local AS by BGP routers, which determine the egress router for the traffic

■ Method

- Selection of the best path, including selection of the egress router, is determined by the **BGP decision process**. Decision is based on **path attributes**.
- The choice of the preferred egress router can be influenced by modifying path attributes for a prefix
- Prefix is manipulated and admitted (propagated) into local AS to let routers decide
- Local AS routers select the path, which is “preferred”, over any other path(s), in accordance with the **BGP decision process**

Example: the use of Local Preference attribute for controlling **outbound** traffic

AS1 view (provider view)


```
Import: from AS2 RX at R1 set localpref=300;  
       from AS2 RX at R2 set localpref=100;  
       accept AS2
```

```
Export: to AS2 RX at R1 announce ANY  
       to AS2 RX at R2 announce ANY
```

AS2 view (customer view)


```
Import: from AS1 R1 at RX set localpref=200  
       from AS1 R2 at RX set localpref=100;  
       accept ANY
```

```
Export: to AS1 R1 at RX announce AS2  
       to AS1 R2 at RX announce AS2
```

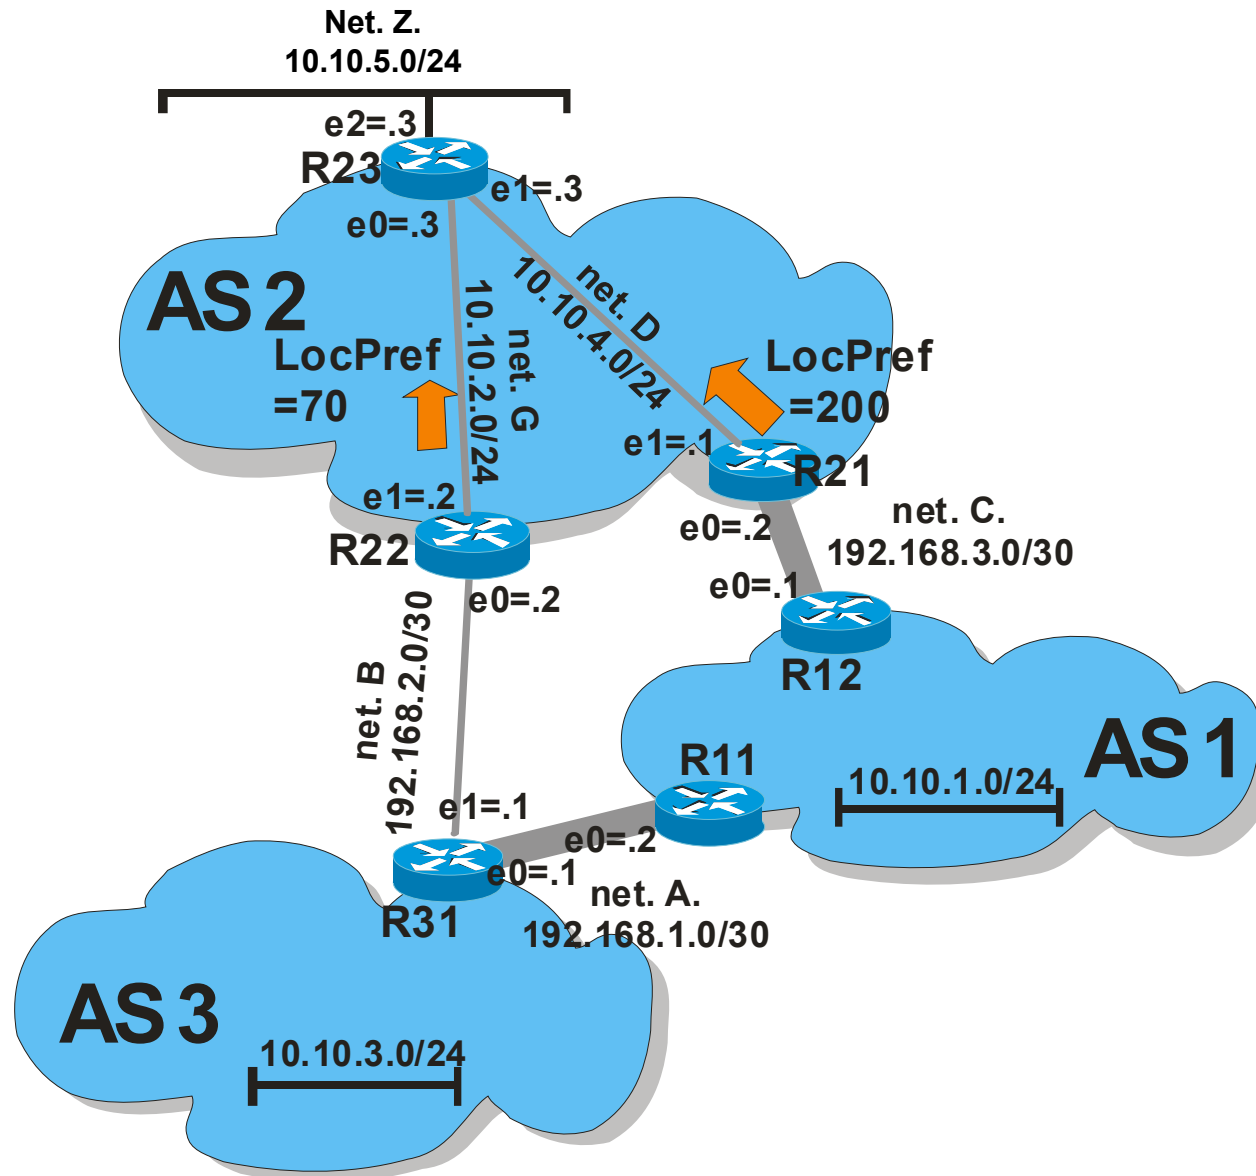


Example for Local Preference

- **LocPref** is set in **AS2** on the received prefix **10.10.1.0/24** (the prefix is received from **AS1** directly and from **AS1** via **AS3**)
- The **LocPref** values indicate that BGP routers in **AS2** "should prefer" path via **AS3** to network **10.10.1.0/24**



Network Topology



Lab Scenario Personalization

- Modify the default scenario in the following way:
 - change the network Z IP address to **<LAB-ID>.10.5.0/24**, where LAB-ID is your personal ID assigned by the lab instructor
- **Note well:** from now-on
 - Command-line commands should reflect this change, therefore there can be differences in the outputs shown in the manual

AS2 (Provider) Perspective

- **Start the lab.** Zebra is 'up and running', interfaces are up, networks are configured.
- Note the networks
 - 10.10.1.0/24 in AS1
 - 10.10.3.0/24 in AS3
- ... but first check configuration of R31, R21, and R22

Check Configuration of R31, R21, and R22

- Step 1
 - Check the general configuration of the **zebra** routing daemon
- Step 2
 - Check if the **bgpd** routing daemon is turned on in **zebra** on R31, R21, and R22
- Step 3
 - Verify if **bgpd** is properly configured

Changing R21 and R22 Configuration

- The **next-hop-self** attribute can be used on R21 and R22 to provide reachability of the next-hop routers (instead of IGP)
 - `#neighbor <peer> next-hop-self`
- „This command specifies an announced route's nexthop as being equivalent to the address of the bgp router” [Quagga manual]
- Additional information on **next-hop-self**:
<http://www.getnetworking.net/bgp/bgp-next-hop-self>

Reachability over BGP

- Run `ping 10.10.3.1` on R23
 - Q1: Which path is selected?
 - Q2: Why this path is chosen and not the other one?
- Q3: Is there any method to make the router R23 choose the other path? What are the possibilities?

Changing R21 and R22 Configuration

- To configure R21 and R22 (see next slide) do **one** of the following
 - Telnet to `bgpd`
 - Configure BGP
 - Issue the `#clear ip bgp *` command
 - Use `#vtysh`
 - Configure BGP
 - Issue the `#clear ip bgp *` command
 - Edit the `bgpd.conf` file
 - Put BGP configuration
 - Restart `zebra`

Exemplary Configuration

```
router bgp xxx
    network nn.nn.nn.nn/mm
    neighbor aa.bb.cc.dd remote-as nnnn
    neighbor aa.bb.cc.dd route-map myRouteMap in
!
route-map myRouteMap permit 10
    set local-preference xx
!
```

Route Map – command syntax

command syntax

```
neighbor <neighbor-ip> route-map <r-map-name> in
```

command syntax

```
neighbor <neighbor-ip> route-map <r-map-name> out
```

command syntax

```
route-map <r-map-name> permit <seq-number>  
  match <announce-property>  
  set <attribute-setting>  
  ...
```

command syntax

```
route-map <r-map-name> deny <seq-number>  
  match <announce-property>  
  set <attribute-setting>  
  ...
```

Testing routes from R23 to AS3

- Execute `ping` from R23 to 10.10.3.1
- Q4: Which path is selected for "ICMP echo request" and "ICMP echo reply" packets?
Why?

Reporting

- Please deliver the following items to the UPEL system using your account
 1. A photocopy or a screenshot showing the output of the following command executed on router **R23**:
 - `ping -I <Lab-ID>.10.5.3 -R 10.10.3.1`